

Nephtys: Lightweight Edge Connector for Bandwidth-Efficient Ingestion of Urban Sensor Streams

Andrea Bozzo
Independent Researcher
andreabozzo92@gmail.com

Abstract—Smart City sensor networks generate continuous telemetry that must be ingested, filtered, and forwarded by resource-conscious edge gateways. Existing tools span forwarding-oriented brokers and broader integration or stream-processing platforms, leaving different trade-offs in processing scope, reconfiguration semantics, and deployment footprint. We present Nephtys, a Go-based edge connector that ingests heterogeneous real-time sources (WebSocket, REST, SSE, webhooks, and gRPC) through per-stream middleware pipelines—composable chains of transformation, deduplication, threshold filtering, and batching stages that can be hot-swapped through a declarative REST API. NATS JetStream provides transport, event durability, and configuration storage without a separate application database. Across three trials with 20 simulated stations and five live feeds, Nephtys reduced bytes by $67.30 \pm 0.03\%$ and $88.96 \pm 0.34\%$, respectively, while using 25.60 ± 0.26 MB connector RSS.

Index Terms—Edge computing, IoT, Smart City, data ingestion, stream processing, NATS, middleware pipeline

I. INTRODUCTION

Smart City deployments and the Industrial Internet of Things (IIoT) have shifted data acquisition from static, batch-oriented collection to continuous streaming from ubiquitous sensor networks [1]. Urban air-quality stations, traffic sensors, and environmental probes can generate redundant telemetry across bandwidth-limited links between edge gateways and centralized backends. General-purpose distributed brokers and stream processors provide broad capabilities, but their operational model may be disproportionate for a gateway whose task is narrowly scoped to acquiring, reducing, and durably forwarding sensor streams [2], [3].

Forwarding mechanisms such as MQTT bridges [4], flow-based tools such as Node-RED [5], and configurable stream processors such as Redpanda Connect [6] cover overlapping parts of this task. Few lightweight designs combine multi-protocol acquisition, per-stream reduction, durable configuration without a separate database, and replacement of processing logic while leaving the source connection running. We investigate this combination as a focused edge-connector design point rather than as a replacement for general IoT or stream-processing platforms.

This paper presents **Nephtys**, a telemetry ingestion connector developed in Go [7] and designed for resource-conscious edge deployment. It ingests WebSocket, REST-polling, Server-Sent Events, webhook, and gRPC sources. Each stream passes through a per-stream **middleware pipeline**: a composable chain of filter, transform, deduplicate, threshold, enrich, and batch stages that can be hot-swapped through a declarative REST API without restarting the source connector. Processed events are published to NATS JetStream [8]; initial stream registrations are stored in a JetStream Key-Value bucket to recover connectors after process restarts without a separate application database.

The contributions of this work are as follows:

- 1) A **multi-protocol ingestion architecture** with native fault tolerance (exponential-backoff reconnection, graceful shutdown) designed for intermittent edge connectivity.
- 2) A **dynamically reconfigurable middleware pipeline** engine that reduces upstream bandwidth by filtering, deduplicating, and compacting sensor payloads at the edge, evaluated on both synthetic and live urban air-quality streams.
- 3) A **database-free configuration persistence** model that uses the existing NATS JetStream subsystem for event durability and initial stream-registration state.

II. SYSTEM ARCHITECTURE AND RESILIENCE

Nephtys targets the constraints of edge computing: limited memory, unreliable network connectivity, and the absence of orchestration infrastructure [2]. Fig. 1 illustrates the overall architecture: urban sensor stations feed data into Nephtys through protocol-specific connectors; each stream traverses a configurable middleware pipeline before publication to NATS JetStream [8], from which cloud-side consumers subscribe.

A. Fault-Tolerant Ingestion

Each data source is managed by a *connector* abstraction that encapsulates the protocol-specific logic (WebSocket, REST poller, SSE, webhook, gRPC). In a Smart City deployment [1], connectors may simultaneously ingest high-




Fig. 1. Nephtys system architecture. Sensor data flows left-to-right through protocol connectors and a composable middleware pipeline before publication to NATS JetStream. The REST API enables runtime hot-swapping of pipeline configurations.

frequency air-quality readings via WebSocket from a local sensor gateway and poll a municipal open-data REST API for meteorological context.

The pull-based WebSocket and SSE connectors reconnect with **exponential backoff** (initial delay 1 s, capped at 30 s), limiting retry pressure when an upstream gateway is unavailable. REST polling retries on its next configured interval. Inbound webhook and gRPC sources instead delegate delivery retry to their upstream clients because they accept rather than initiate connections.

B. Database-Free Configuration Persistence

To avoid operating a separate application database, Nephtys uses **NATS JetStream** [8] as both event transport and configuration store. Published sensor events are retained in a durable JetStream stream (72 hours by default). Initial stream registrations are serialized in a JetStream Key-Value bucket (`nephtys_streams`); on startup, the Stream Manager reads these entries and re-registers their connectors and pipelines. Pipeline changes made through the hot-swap endpoint are intentionally transient in the evaluated version and revert to the registered configuration after restart.

This design keeps the required application stack to Nephtys and NATS; Prometheus and Grafana are optional evaluation tools. Outbound WebSocket and HTTP-based connectors accept TLS-protected endpoints, and the management REST API supports Bearer-token authorization. Actual single-board-computer measurements remain future validation rather than an assumption of the present experiment.

C. Edge Pipeline for Bandwidth Reduction

The central component is the **Pipeline Middleware** engine. Before an event reaches the broker, it traverses a chain of composable middleware stages, each independently configurable per stream:

- **Transform:** Extracts only the relevant fields from nested JSON payloads using dot-notation path mapping, reducing per-event payload size.
- **Dedup:** Computes the FNV-64a hash [9] of each payload and blocks duplicates within a configurable time-to-live window.
- **Threshold:** Passes an event only if the absolute change in a numerical field exceeds a configurable delta (e.g., $\Delta \text{PM}_{2.5} \geq 1.0 \mu\text{g}/\text{m}^3$), filtering out sensor noise.
- **Batch:** Buffers events and flushes them as a single aggregated message at a configurable interval or batch size, reducing per-message overhead.
- **Filter and Enrich:** Drop events not matching specified types, or inject static metadata tags.

The pipeline configuration is a JSON object attached to each stream registration and can be **hot-swapped at runtime** via a `PUT /v1/streams/{id}/pipeline` REST call. Internally, the active pipeline handler is stored behind a Go `atomic.Pointer`, so the swap is lock-free and takes effect on the next event with sub-millisecond latency. This allows operators to tighten or relax filtering thresholds in response to air-quality alerts without interrupting data flow.

III. EVALUATION SCENARIO

We evaluate Nephtys in an urban air-quality monitoring scenario representative of a Smart City edge deployment. All source code, the sensor simulator, and an automated benchmark script are publicly available¹ to allow full reproducibility of the results presented below.

A. Experimental Setup

Nephtys runs natively on Windows 11 (Intel Core Ultra 7 258V, 32 GB RAM); NATS JetStream, Prometheus, and Grafana are co-deployed via Docker Compose. The

¹<https://github.com/AndreaBozzo/uic2026-nephtys> (companion material); <https://github.com/AndreaBozzo/Nephtys> (Nephtys source).

TABLE I
BENCHMARK RESULTS OVER THREE TRIALS (MEAN; $\pm$ SAMPLE SD).

Metric	Baseline	Pipeline	Reduction
<i>Synthetic streams (20 stations, 5 min)</i>			
Bytes published	2,736,478	894,897	67.30 $\pm$ 0.03%
Messages published	12,027	155	98.71 $\pm$ 0.00%
<i>Real-data streams (5 APIs, 30 min)</i>			
Bytes published	142,489	15,729	88.96 $\pm$ 0.34%
Messages published	270	20	92.71 $\pm$ 0.22%
RSS memory	25.60 MB		$\pm$ 0.26 MB
RSS via Prometheus <code>process_resident_memory_bytes</code> at steady state.			

evaluated Nephtys revision is `c146ee7`. Two classes of data sources feed the system:

- 1) **Synthetic streams.** A lightweight WebSocket-based simulator (`sensor-sim`) generates air-quality readings for 20 virtual stations at 2Hz, producing $\sim$ 40 events/s. Each event carries PM_{2.5}, PM₁₀, NO₂, temperature, and humidity values with configurable jitter and a 30% duplicate ratio, modeling the redundancy typical of low-cost particulate-matter sensors.
- 2) **Real-world open data.** Five REST poller streams ingest live measurements from the Open-Meteo weather and air-quality APIs [10] for three cities in Calabria (Cosenza, Catanzaro, Reggio Calabria) and from a Sensor.Community [11] citizen PM sensor near Gioia Tauro, polled every 30–60 seconds.

The full middleware pipeline is configured as: Transform (5 fields extracted), Dedup (TTL = 30s, cache = 500), Threshold (path = `pm25`, $\delta = 1.0 \mu\text{g}/\text{m}^3$), Batch (size = 50, flush = 5s).

B. Results

Table I reports three trials of two experiments: (a) a 5-minute synthetic workload with 20 virtual stations at 2 Hz ($\sim$ 40 events/s, 30% duplicate ratio), and (b) a 30-minute live workload with five REST poller streams hitting the Open-Meteo [10] and Sensor.Community [11] APIs. The simulator uses a fixed seed. Each table value is the arithmetic mean; reduction ratios and RSS additionally show sample standard deviation. Live inputs vary with the upstream APIs.

Synthetic workload. Per trial, the pipeline ingested a mean of 12,027 events. **Dedup** dropped 3,574 and **Threshold** another 686 on average; the remaining events were compacted by **Batch** into 155 messages (about 50 events per batch). Together with **Transform**, this reduced published bytes by 67.30% and message count by 98.71%.

Real-data workload. Across the five live API streams, each 30-minute trial ingested 270 events on average; Dedup dropped 158 and Threshold 92, leaving 20 published messages. Transform and filtering jointly reduced bytes by 88.96% and message count by 92.71%. One trial experi-




Fig. 2. Grafana dashboard during a mixed workload (synthetic + real-data streams), showing ingested vs. published rates, per-middleware drop breakdown, and instantaneous reduction gauges (65.7% bandwidth, 98.6% events at capture time; final aggregates in Table I).

TABLE II
QUALITATIVE COMPARISON OF OVERLAPPING EDGE-INGESTION TOOLS.

System	Primary scope	Runtime change	Config state
MQTT bridge	Broker forwarding	Reconnect/reload	Broker-specific
Node-RED	General visual flows	Flow redeploy	Local flow file
EdgeX	IoT microservices	Service/rule APIs	Service databases
Redpanda Connect	General pipelines	Stream replacement	Files; runtime API
Nephtys	Source-to-NATS edge	Pipeline-only swap	KV registration

enced a Sensor.Community connection reset and recovered at the next poll, illustrating live-input variability.

Resilience. When the sensor simulator process is killed, the WebSocket connector detects the disconnection and enters exponential backoff (1s, 2s, 4s, ..., capped at 30s). After the simulator is restarted, Nephtys reconnects automatically within 1–2s. No crash or manual intervention is required; ingestion simply pauses and resumes.

Hot-reload latency. Issuing a PUT request to change the threshold delta completes in under 16ms (including the full HTTP round-trip). The atomic pointer swap ensures the new pipeline takes effect on the very next event, with no loss or reordering.

Fig. 2 shows the Grafana dashboard captured during the benchmark, displaying the real-time bandwidth and event reduction ratios.

IV. RELATED WORK AND CONCLUSIONS

A. Related Work

Table II compares capabilities rather than mixing memory figures measured under different workloads. EdgeX Foundry [12] is a broad microservice-based IoT platform, while MQTT bridges [4] focus on broker-to-broker forwarding. Node-RED [5] provides general flow-based processing and persists deployed flows to files. Redpanda Connect is the closest configuration-driven comparator: its streams mode exposes runtime CRUD for isolated processor pipelines, but an update shuts down and replaces the affected stream [6]. Nephtys deliberately offers fewer processors and sinks; its narrower distinction is swapping a processing chain while retaining the running source connector, coupled with JetStream-backed registration state.

B. Conclusions

This paper presented Nephtys, a lightweight edge connector that demonstrates how compiled languages and lightweight message brokers can enable bandwidth-efficient telemetry ingestion for Smart City sensor networks. By shifting deduplication, transformation, and threshold filtering to the edge through declaratively configurable middleware pipelines [3], the system reduced bytes by 67.30% on the synthetic workload and 88.96% on the live workload, with 25.60 MB mean connector RSS.

The runtime-reconfigurable pipeline allows operators to adapt filtering behavior (e.g., tightening thresholds during pollution alerts) without restarting the source connector. The contribution is the evaluated combination of this update boundary, heterogeneous acquisition, and reuse of the NATS durability layer—not a claim that runtime-reconfigurable stream processing is unique to Nephtys.

Current limitations include only three trials, a 30-minute live-data window with changing upstream values, host-machine rather than single-board-computer measurements, no controlled quantitative platform baseline, and a single-node deployment. The evaluated system also uses at-least-once delivery, and pipeline hot-swaps are not persisted across restart. These constraints limit statistical and platform-level generalization of the reported results.

Immediate future work is a repeated, controlled evaluation on edge hardware against a representative flow-based baseline. Longer-term work will investigate adaptive pipeline policies that trade information loss and processing latency against bandwidth and resource budgets, using

multi-node sensor deployments only after the single-node operating envelope is characterized.

REFERENCES

- [1] A. Zanella, N. Bui, A. Castellani, L. Vangelista, and M. Zorzi, “Internet of things for smart cities,” *IEEE Internet of Things Journal*, vol. 1, no. 1, pp. 22–32, 2014.
- [2] W. Shi, J. Cao, Q. Zhang, Y. Li, and L. Xu, “Edge computing: Vision and challenges,” *IEEE Internet of Things Journal*, vol. 3, no. 5, pp. 637–646, 2016.
- [3] M. D. de Assunção, A. da Silva Veith, and R. Buyya, “Distributed data stream processing and edge computing: A survey on resource elasticity and future directions,” *Journal of Network and Computer Applications*, vol. 103, pp. 1–17, 2018.
- [4] A. Banks, E. Briggs, K. Borgendale, and R. Gupta, “MQTT version 5.0,” OASIS, OASIS Standard, 2019, 07 March 2019. [Online]. Available: <https://docs.oasis-open.org/mqtt/mqtt/v5.0/os/mqtt-v5.0-os.html>
- [5] OpenJS Foundation, “Node-RED – low-code programming for event-driven applications,” <https://nodered.org>, 2026, accessed: 2026-03-26.
- [6] Redpanda Data, “Redpanda Connect streams api,” https://docs.redpanda.com/connect/guides/streams_mode/streams_api/, 2026, accessed: 2026-07-13.
- [7] The Go Authors, “The Go programming language specification,” <https://go.dev/ref/spec>, 2026, accessed: 2026-03-26.
- [8] NATS Authors, “NATS – the cloud native messaging system,” <https://nats.io>, 2026, accessed: 2026-03-26.
- [9] G. Fowler, L. C. Noll, K.-P. Vo, D. E. Eastlake, III, and T. Hansen, “The FNV non-cryptographic hash algorithm,” IETF Internet-Draft, <https://datatracker.ietf.org/doc/html/draft-eastlake-fnv-17>, 2019, internet-Draft (expired).
- [10] P. Zippenfenig, “Open-Meteo – free weather API,” <https://open-meteo.com>, 2026, accessed: 2026-03-26.
- [11] Sensor.Community, “Sensor.Community – a contributors driven global sensor network,” <https://sensor.community>, 2026, accessed: 2026-03-26.
- [12] LF Edge, “EdgeX Foundry – open-source edge computing platform for IoT,” <https://www.edgexfoundry.org>, 2026, accessed: 2026-03-26.